

# 11. R中的并行计算III

罗翔宇

中国人民大学统计与大数据研究院

此课件内容基于Norman Matloff编著的《数据科学中的并行计算》（CRC Press, 汪磊、寇强译）

# 上周复习

- ▶ 基于共享内存范式的Rdsm包
- ▶ myinfo\$nwkrks 线程个数, myinfo\$id 当前线程id
- ▶ mgrinit()初始化Rdsm包系统
- ▶ mgrmakevar()在共享内存中创建变量
- ▶ 用snow包中的clusterEvalQ()来开始线程
- ▶ 锁: 通过锁变量, 将读和写看成一个不可分割的组合, 即原子化
- ▶ 屏障: 用来同步所有线程

# 例子1：时间序列中的最大脉冲

- ▶ 考虑一个长度为 $n$ 的时间序列，我们可能对脉冲感(一个时间段)感兴趣，也就是一段时间内保持较高的平均值。我们希望找到长度为 $k$ 的具有最大平均值的周期
- ▶ 时间序列zoo包中有一个函数`rollmean(w,k)`,它返回所有长度为 $k$ 的块各个平均值

# 例子1：时间序列中的最大脉冲

```
maxburst <- function(x, k, max, rsIts){  
  require(Rdsm)  
  require(zoo)  
  
  n <- length(x)  
  myidxs <- getidxs(n-k+1)  
  myfirst <- myidxs[1]  
  mylast <- myidxs[length(myidxs)]  
  mas[1, myfirst:mylast] <- rollmean(x[myfirst:(mylast+k-1)], k)  
  
  barr()  
  
  if(myinfo$id == 1){  
    rsIts[1,1] <- which.max(mas[1, ])  
    rsIts[1,2] <- mas[1, rsIts[1,1]]  
  }  
}
```

- ▶ `getidxs(n-k+1)`将1到n-k+1根据线程个数分为若干段
- ▶ 通过`rollmean(x[myfirst:(mylast+k-1)], k)`将此段中的周期k的脉冲的平均值计算出来（即长度为k的脉冲的开头分别从myfirst到mylast）
- ▶ `barr()`确保所有线程都执行完了，指派一个线程合并各个进程的结果

# 例子1：时间序列中的最大脉冲

## ► 运行所得结果

```
test <- function(cls){  
  require(Rdsm)  
  mgrinit(cls)  
  mgrmakevar(cls, "mas", 1, 9)  
  mgrmakevar(cls, "rslts", 1, 2)  
  x <- c(5, 7, 6, 20, 4, 14, 11, 12, 15, 17)  
  clusterExport(cls, "maxburst")  
  clusterExport(cls, "x")  
  clusterEvalQ(cls, maxburst(x, 2, mas, rslts))  
  print(rslts[,])  
}  
  
c2 <- makeCluster(2)  
test(c2)
```

```
> c2 <- makeCluster(2)  
> test(c2)  
[1]  9 16
```

## 例子2：变换邻接矩阵

- ▶ 假设我们有一个图，它的邻接矩阵是

0   1   1

- ▶ 0   0   0

1   1   0

- ▶ 其中，第1行第2列为1，表示顶点1到顶点2有边；第1行第3列为1，表示顶点1到顶点3有边。

- ▶ 所对应的图为


- ▶ 我们的目标为将其变成一个两列矩阵

1,2

1,3

3,1

3,2



## 例子2：变换邻接矩阵

### ► 代码：

```
convert1row <- function(rownum, colswith1s){  
  if(is.null(colswith1s)) return(NULL)  
  cbind(rownum, colswith1s)  
}
```

### ► 将某一行的所得1的位置，转化为矩阵形式

### ► 比如，第3行，在(2, 5, 9)的位置为1，那么返回的结果为

3,2

3,5

3,9



## 例子2: 变换邻接矩阵

```
findlinks <- function(adj, lnks, counts){  
  nr <- nrow(adj)  
  
  myidxs <- getidxs(nr)  
  
  myout <- apply(adj[myidxs, ], 1, function(onerow) which(onerow==1))  
  
  tmp <- matrix(nrow=0,ncol=2)  
  my1strow <- myidxs[1]  
  
  for(idx in myidxs)  
    tmp <- rbind(tmp, convert1row(idx, myout[[idx-my1strow+1]]))  
  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  me <- myinfo$id  
  counts[1, me] <- nmyedges  
  barr()  
  
  if(me == 1) counts[1, ] <- cumsum(counts[1, ])  
  barr()  
  
  mystart <- if(me == 1) 1 else counts[1, me-1]+1  
  myend <- mystart + nmyedges - 1  
  lnks[mystart:myend, ] <- tmp  
  
  0  
}
```

- 输入中, adj表示邻接矩阵( $n$  by  $n$ ), lnks表示存储的矩阵( $n^2$  by 2), counts表示每个线程得到的边的数目

## 例子2：变换邻接矩阵

```
findlinks <- function(adj, lnks, counts){  
  nr <- nrow(adj)  
  
  myidxs <- getidxs(nr)  
  
  myout <- apply(adj[myidxs, ], 1, function(onerow) which(onerow==1))  
  
  tmp <- matrix(nrow=0,ncol=2)  
  my1strow <- myidxs[1]  
  
  for(idx in myidxs)  
    tmp <- rbind(tmp, convert1row(idx, myout[[idx-my1strow+1]]))  
  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  me <- myinfo$id  
  counts[1, me] <- nmyedges  
  barr()  
  
  if(me == 1) counts[1, ] <- cumsum(counts[1, ])  
  barr()  
  
  mystart <- if(me == 1) 1 else counts[1, me-1]+1  
  myend <- mystart + nmyedges - 1  
  lnks[mystart:myend, ] <- tmp  
  
  0  
}
```

- getidxs将邻接矩阵的行分为若干块，每个线程通过apply()处理对应块中边的统计

## 例子2：变换邻接矩阵

```
findlinks <- function(adj, lnks, counts){  
  nr <- nrow(adj)  
  
  myidxs <- getidxs(nr)  
  
  myout <- apply(adj[myidxs, ], 1, function(onerow) which(onerow==1))  
  
  tmp <- matrix(nrow=0,ncol=2)  
  my1strow <- myidxs[1]  
  
  for(idx in myidxs)  
    tmp <- rbind(tmp, convert1row(idx, myout[[idx-my1strow+1]]))  
  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  me <- myinfo$id  
  counts[1, me] <- nmyedges  
  barr()  
  
  if(me == 1) counts[1, ] <- cumsum(counts[1, ])  
  barr()  
  
  mystart <- if(me == 1) 1 else counts[1, me-1]+1  
  myend <- mystart + nmyedges - 1  
  lnks[mystart:myend, ] <- tmp  
  
  0  
}
```

- 在for循环中，每个线程遍历对应块中的每一行，利用convert1row将边的位置转化为一个矩阵，最后组成一个总的矩阵

## 例子2：变换邻接矩阵

```
findlinks <- function(adj, lnks, counts){  
  nr <- nrow(adj)  
  
  myidxs <- getidxs(nr)  
  
  myout <- apply(adj[myidxs, ], 1, function(onerow) which(onerow==1))  
  
  tmp <- matrix(nrow=0,ncol=2)  
  my1strow <- myidxs[1]  
  
  for(idx in myidxs)  
    tmp <- rbind(tmp, convert1row(idx, myout[[idx-my1strow+1]]))  
  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  me <- myinfo$id  
  counts[1, me] <- nmyedges  
  barr()  
  
  if(me == 1) counts[1, ] <- cumsum(counts[1, ])  
  barr()  
  
  mystart <- if(me == 1) 1 else counts[1, me-1]+1  
  myend <- mystart + nmyedges - 1  
  lnks[mystart:myend, ] <- tmp  
  
  0  
}
```

- nmyedges为每个线程所找到的边的个数，将此信息放入counts中。为了将结果汇总在lnks中，通过在线程1中对counts累加，得到mystart

## 例子2：变换邻接矩阵

```
findlinks <- function(adj, lnks, counts){  
  nr <- nrow(adj)  
  
  myidxs <- getidxs(nr)  
  
  myout <- apply(adj[myidxs, ], 1, function(onerow) which(onerow==1))  
  
  tmp <- matrix(nrow=0,ncol=2)  
  my1strow <- myidxs[1]  
  
  for(idx in myidxs)  
    tmp <- rbind(tmp, convert1row(idx, myout[[idx-my1strow+1]]))  
  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  me <- myinfo$id  
  counts[1, me] <- nmyedges  
  barr()  
  
  if(me == 1) counts[1, ] <- cumsum(counts[1, ])  
  barr()  
  
  mystart <- if(me == 1) 1 else counts[1, me-1]+1  
  myend <- mystart + nmyedges - 1  
  lnks[mystart:myend, ] <- tmp  
  
  0  
}
```

- 通过mystart和myend信息，将每个线程对应的结果存入lnks的对应列中

## 例子2：变换邻接矩阵

```
test <- function(adj, n, cls){  
  mgrinit(cls)  
  mgrmakevar(cls, "x", n, n)  
  mgrmakevar(cls, "lnks", n^2, 2)  
  mgrmakevar(cls, "counts", 1, length(cls))  
  
  x[, ] <- adj  
  clusterExport(cls, "findlinks")  
  clusterExport(cls, "convert1row")  
  clusterEvalQ(cls, findlinks(x, lnks, counts))  
  print(lnks [1:counts[1, length(cls)], ])  
}
```

- 做一个测试，adj为邻接矩阵，n为顶点个数，cls为分配的snow集群

# 例子2：变换邻接矩阵

## ► 结果

```
library(parallel)
library(Rdsm)
```

```
n <- 5
adj <- matrix(sample(0:1, n^2, replace = TRUE), ncol = n)
lnks <- matrix(NA, n^2, 2)
cls <- makeCluster(2)
test(adj, n, cls)
```

```
> adj
      [,1] [,2] [,3] [,4] [,5]
[1,]    0    0    0    0    1
[2,]    1    0    0    1    1
[3,]    1    0    1    1    1
[4,]    1    0    0    1    1
[5,]    0    0    1    0    0
```

```
> test(adj, n, cls)
      [,1] [,2]
[1,]    1    5
[2,]    2    1
[3,]    2    4
[4,]    2    5
[5,]    3    1
[6,]    3    3
[7,]    3    4
[8,]    3    5
[9,]    4    1
[10,]   4    4
[11,]   4    5
[12,]   5    3
```

# 例子2：变换邻接矩阵

## ► 串行代码

```
getlinksnonpar <- function(adj, lns){  
  nr <- nrow(adj)  
  myout <- apply(adj, 1, function(rw) which(rw==1))  
  nmyedges <- Reduce(sum, lapply(myout, length))  
  lnsidx <- 1  
  
  for(idx in 1:nr){  
    jdx <- idx  
    myoj <- myout[[jdx]]  
    endwrite <- lnsidx + length(myoj) - 1  
    if(!is.null(myoj)){  
      lns[lnsidx:endwrite, ] <- cbind(idx, myoj)  
    }  
    lnsidx <- endwrite + 1  
  }  
}
```



# 例子2：变换邻接矩阵

## ► 串行代码

```
n <- 10000  
adj <- matrix(sample(0:1, n^2, replace = TRUE), ncol = n)  
lnks <- matrix(NA, n^2, 2)  
system.time(getlinksnonpar(adj, lnks))
```

```
> system.time(getlinksnonpar(adj, lnks))  
   user  system elapsed  
8.373   0.492   8.867
```

# 例子2：变换邻接矩阵

## ► 并行代码

```
cls <- makecluster(8)
mgrinit(cls)
mgrmakevar(cls, "x", n, n)
mgrmakevar(cls, "lnks", n^2, 2)
mgrmakevar(cls, "counts", 1, length(cls))

x[, ] <- adj
clusterExport(cls, "findlinks")
clusterExport(cls, "convert1row")
system.time(clusterEvalQ(cls, findlinks(x, lnks, counts)))
```

```
user  system elapsed
0.002   0.000   53.710
```

## ► 可能原因为某个线程耽误了太多时间，导致其它线程的长时间等候

## 例子2：变换邻接矩阵

- ▶ 为了给矩阵lnks分配内存来处理最坏的情形，我们浪费了空间和执行时间。问题在于我们并不能提前知道“输出”的大小，也就是lnks的大小
- ▶ 在测试代码中，用户通过检查counts[1, length(cls)]来确定lnks中“真实”的行数，把这些“真实”的行复制给另一个矩阵，然后释放大矩阵的内存

# 例子3: kmeans聚类

- ▶ kmeans将数据聚成k个组，现在我们通过Rdsm包进行实现
- ▶ 不断迭代以下步骤
- ▶ 1.对每个数据点，即数据矩阵中的每一行，确定这个点离哪个中心点最近
- ▶ 2.把这个数据点添加到那个中心所在的组内
- ▶ 3.更新中心

# 例子3: kmeans聚类

## ► 代码

```
kmeans_rdsm <- function(x, k, ni, cntrds, sums, lck, cinit = NULL){  
  require(parallel)  
  require(pdist)  
  nx <- nrow(x)  
  
  myidxs <- getidxs(nx)  
  myx <- x[myidxs, ]  
  
  if(is.null(cinit)){  
    if(myinfo$id == 1)  
      cntrds[, ] <- x[sample(1:nx, k, replace = F), ]  
    barr()  
  }else{  
    cntrds[, ] <- cinit  
  }  
  
  mysum <- function(idxs, myx){  
    c(length(idxs), colSums(myx[idxs, , drop=F]))  
  }  
  
  for(i in 1:ni){  
    if(myinfo$id == 1){  
      sums[] <- 0  
    }  
  
    barr()  
  
    #calculate distances  
    dsts <- matrix(pdist(myx, cntrds[, , drop=F])@dist, ncol = nrow(myx))  
  
    nrst <- apply(dsts, 2, which.min)
```

# 例子3: kmeans聚类

- ▶  $X$ 表示数据矩阵，行为样本，列为维度； $k$ 表示聚类数目； $ni$ 表示迭代次数； $cntrds$ 表示中心点矩阵，行为类别，列为维度； $sums$ 表示一个矩阵，行为类别，第一列表示对应类别有多少样本量，其余列表示在不同维度上相同类别值的和； $lck$ 表示锁变量； $cinit$ 表示是否初始化，默认为空
- ▶  $getidxs$ 每个线程分配到对应块； $myx$ 表示线程的对应块
- ▶ 若 $cinit$ 为空，就用随机获得 $cntrds$ 初始值，由于是共享变量，只需要在线程1中赋值即可。
- ▶ 利用 $barr()$ 让所有线程等待线程1的赋值
- ▶  $mysum$ 是一个自定义函数。其最终目的是利用 $tapply$ 函数返回线程所负责那块数据的结果总结。该函数返回为一个向量，第一个元表示对应指标的个数，其余元表示不同维度上值的求和

# 例子3: kmeans聚类

```
#update centers
tmp <- tapply(1:nrow(myx), nrst, mysum, myx)

realrdsmlock(lck)
for(j in as.integer(names(tmp))){
  sums[j, ] <- sums[j, ] + tmp[[j]]
}
realrdsmunlock(lck)
barr()

if(myinfo$id == 1){
  for(j in 1:k){
    if(sums[j, 1] > 0){
      cntrds[j, ] <- sums[j, -1] / sums[j, 1]
    }else{
      cntrds[j, ] <- x[sample(1:nx, 1), ]
    }
  }
}

}
```

- ▶ `tapply`返回的是一个列表，列表的部分1表示线程负责的此块数据中属于类别1的信息，第一元表示该类别中的数据个数，其它元表示该类别的数据在不同维度上的求和
- ▶ `realrdsmlock()`通过锁操作将不同线程的结果汇总在sums，可以避免线程之间的竞争导致错误的结果

# 例子3: kmeans聚类

## ► 一个测试函数

```
test <- function(cls){  
  mgrinit(cls)  
  mgrmakevar(cls, "x", 6, 2)  
  mgrmakevar(cls, "cntrds", 2, 2)  
  mgrmakevar(cls, "sms", 2, 3)  
  mgrmakelock(cls, "lck")  
  
  x[, ] <- matrix(sample(1:20, 12), ncol=2)  
  clusterExport(cls, "kmeans_rdsm")  
  clusterExport(cls, "pdist")  
  clusterEvalQ(cls, kmeans_rdsm(x, 2, 1, cntrds, sms,  
                                "lck", cinit=rbind(c(5, 5), c(15, 15))))  
  cntrds[, ]  
}
```



# 例子3: kmeans聚类

## ► 另一个测试函数

```
test1 <- function(cls){  
  mgrinit(cls)  
  mgrmakevar(cls, "x", 10000, 3)  
  mgrmakevar(cls, "cntrds", 3, 3)  
  mgrmakevar(cls, "sms", 3, 4)  
  mgrmakelock(cls, "lck")  
  
  x[, ] <- matrix(rnorm(30000), ncol=3)  
  ri <- sample(1:10000, 3000)  
  x[ri, 1] <- x[ri, 1] + 5  
  ri <- sample(1:10000, 3000)  
  x[ri, 2] <- x[ri, 2] + 5  
  
  clusterExport(cls, "kmeans_rdsm")  
  clusterExport(cls, "pdist")  
  clusterEvalQ(cls, kmeans_rdsm(x, 3, 50, cntrds, sms, "lck"))  
  cntrds[,]  
}
```

# 例子3: kmeans聚类

## ► 结果

```
library(parallel)
library(Rdsm)
library(pdist)

cls <- makeCluster(2)

test(cls)

system.time(test1(cls))

stopCluster(cls)
```

```
> test(cls)
      [,1] [,2]
[1,]  7.50  5.5
[2,] 12.75 10.5
```

```
> system.time(test1(cls))
   user  system elapsed 
0.049   0.002   1.600
```

# 共享内存范式：C语言层面

- ▶ 直接在多核机器上编程的标准方法是使用线程库，例如在Unix家族的系统上，pthreads库是相当流行的
- ▶ 人们开发了专门考虑并行计算的高级库，包括OpenMP, Intel的Threads Building Blocks 以及 Cilk++
- ▶ 在这里我们简单介绍下OpenMP，以下内容根据<https://helloacm.com/simple-tutorial-with-openmp-how-to-use-parallel-block-in-c-using-openmp/>

# Hello OpenMp


```
#include <stdio.h>

int main(){
    #pragma omp parallel
    {
        printf("hellow openmp!\n");
    }
    return 0;
}
```

► 编译命令为 `gcc -o file file.c -fopenmp`

```
[luoxiangyu@rmdx-cluster openmp_files]$ gcc -o hello_openmp hello_openmp.c -fopenmp
[luoxiangyu@rmdx-cluster openmp_files]$ ./hello_openmp
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
```

# Hello OpenMp



- ▶ `#pragma omp parallel` 用来声明 “从下面开始，我要用多个线程同时执行任务了~”
- ▶ 根据CPU的核数，一些线程将会生成用来执行任务
- ▶ 其实我们也可以设置线程个数

# 设置线程个数 num\_threads()

```
#include <stdio.h>

int main(){
    int x = 1;
    int y = x + 1;
    #pragma omp parallel num_threads(y*3)
    {
        printf("hellow openmp!\n");
    }
    return 0;
}
```

```
[luoxiangyu@rmdx-cluster openmp_files]$ ./hello_openmp_num_threads
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
hellow openmp!
```

# 获得每个线程的信息

```
#include <stdio.h>
#include <omp.h>

int main(){
    #pragma omp parallel num_threads(3)
    {
        int id = omp_get_thread_num();
        int data = id;
        int total = omp_get_num_threads();
        printf("Greetings from process %d out of %d with Data %d\n",
              id, total, data);
    }
    printf("parallel for ends.\n");
    return 0;
}
```

- ▶ 需要用到头文件 `omp.h`
- ▶ `omp_get_thread_num()` 得到每个线程的id
- ▶ `omp_get_num_threads()` 得到线程的总个数

# 获得每个线程的信息

```
#include <stdio.h>
#include <omp.h>

int main(){
    #pragma omp parallel num_threads(3)
    {
        int id = omp_get_thread_num();
        int data = id;
        int total = omp_get_num_threads();
        printf("Greetings from process %d out of %d with Data %d\n",
              id, total, data);
    }
    printf("parallel for ends.\n");
    return 0;
}
```

```
[luoxiangyu@rmdx-cluster openmp_files]$ ./num_threads
Greetings from process 2 out of 3 with Data 2
Greetings from process 0 out of 3 with Data 0
Greetings from process 1 out of 3 with Data 1
parallel for ends.
```

- ▶ 请注意，id, data, total是在每个线程中定义，不为共享变量



# 设置共享变量以及私有(private)变量

```
#include <stdio.h>
#include <omp.h>

int main(){
    int data;
    #pragma omp parallel num_threads(3)
    {
        int id = omp_get_thread_num();
        data = id;
        int total = omp_get_num_threads();
        printf("Greetings from process %d out of %d with Data %d\n",
              id, total, data);
    }
    printf("parallel for ends.\n");
    return 0;
}
```

```
Greetings from process 2 out of 3 with Data 0
Greetings from process 1 out of 3 with Data 1
Greetings from process 0 out of 3 with Data 0
parallel for ends.
```

- ▶ 在这里data为共享变量。是否可以将共享变量转为私有变量呢

# 设置共享变量以及私有(private)变量

```
#include <stdio.h>
#include <omp.h>

int main(){
    int data, id, total;
    //each thread has its own copy of data, id and total
    #pragma omp parallel private(data, id, total) num_threads(6)
    {
        id = omp_get_thread_num();
        data = id;
        total = omp_get_num_threads();
        printf("Greetings from process %d out of %d with Data %d\n",
              id, total, data);
    }
    printf("parallel for ends.\n");
    return 0;
}
```

- ▶ 通过private()的形式，设置已声明好的变量为私有变量

```
Greetings from process 2 out of 6 with Data 2
Greetings from process 5 out of 6 with Data 5
Greetings from process 0 out of 6 with Data 0
Greetings from process 3 out of 6 with Data 3
Greetings from process 1 out of 6 with Data 1
Greetings from process 4 out of 6 with Data 4
parallel for ends.
```

# 设置临界区域

```
#include <stdio.h>
#include <omp.h>

int main(){
    int data=0;

    #pragma omp parallel num_threads(20)
    {
        data = data + 1;
    }
    printf("data = %d.\n", data);
    return 0;
}
```

```
[luoxiangyu@rmdx-cluster openmp_files]$ ./critical_section
data = 7.
```

得到的结果并不为20

# 设置临界区域

```
#include <stdio.h>
#include <omp.h>

int main(){
    int data=0;

    #pragma omp parallel num_threads(20)
    {
        #pragma omp critical
        {
            data = data + 1;
        }
    }
    printf("data = %d.\n", data);
    return 0;
}
```

```
[luoxiangyu@rmdx-cluster openmp_files]$ ./critical_section
data = 20.
```

设置临界区域，使得openmp在区域前后自动设置锁变量，因此得到的结果为20