

5. MPI集体通信及分组


罗翔宇

中国人民大学统计与大数据研究院

此课件内容主要基于Blaise Barney的网络资料
Message Passing Interface (MPI)及Wes Kendall
的MPI Tutorial


上周课堂复习

- ▶ `int MPI_Scatter(void* sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuff, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`




上周课堂复习

- ▶ `int MPI_Gather(void* sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuff, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm);`



上周课堂复习

- ▶ `int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`



上周课堂复习

- `MPI_MAX` - Returns the maximum element.
- `MPI_MIN` - Returns the minimum element.
- `MPI_SUM` - Sums the elements.
- `MPI_PROD` - Multiplies all elements.
- `MPI LAND` - Performs a logical *and* across the elements.
- `MPI_LOR` - Performs a logical *or* across the elements.
- `MPI_BAND` - Performs a bitwise *and* across the bits of the elements.
- `MPI BOR` - Performs a bitwise *or* across the bits of the elements.
- `MPI_MAXLOC` - Returns the maximum value and the rank of the process that owns it.
- `MPI_MINLOC` - Returns the minimum value and the rank of the process that owns it.

MPI缩减操作例子：计算平均数

- ▶ 进程0有数组{1,2,...,100}
- ▶ 进程1有数组{101,102,...,200}
- ▶ 进程2有数组{201,202,...,300}
- ▶ 求每个位置上的平均数
- ▶ 利用算子：MPI_SUM

MPI缩减操作例子：计算平均数

```
MPI_Init(NULL, NULL);  
int world_rank, world_size;  
MPI_Comm_size(MPI_COMM_WORLD, &world_size);  
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
```

```
float num[100], *num_avg;
```

```
for(int i=0; i<100; i++){  
    num[i] = world_rank*100+i+1;  
}
```

```
if(world_rank == 0){  
    num_avg = (float *)malloc(100*sizeof(float));  
}
```

```
MPI_Reduce(num, num_avg, 100, MPI_FLOAT, MPI_SUM,  
           0, MPI_COMM_WORLD);
```

```
if(world_rank == 0){  
    for(int i=0; i<100; i++){  
        num_avg[i] /= world_size;  
        printf(" %f ", num_avg[i]);  
    }  
    free(num_avg);  
}
```

```
MPI_Finalize();
```

所有进程都需要声明
*num_avg


只在进程0中
分配内存，
以节约内存

MPI缩减操作例子：计算平均数

```
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpirun -np 3 ./avg_reduce
101.000000 102.000000 103.000000 104.000000 105.000000 106.000000 107.000000 108.000000 109.000000 110.000000 111.000000 112.000000 113.000000 114.000000 115.000000 116.000000 117.000000 118.000000 119.000000 120.000000 121.000000 122.000000 123.000000 124.000000 125.000000 126.000000 127.000000 128.000000 129.000000 130.000000 131.000000 132.000000 133.000000 134.000000 135.000000 136.000000 137.000000 138.000000 139.000000 140.000000 141.000000 142.000000 143.000000 144.000000 145.000000 146.000000 147.000000 148.000000 149.000000 150.000000 151.000000 152.000000 153.000000 154.000000 155.000000 156.000000 157.000000 158.000000 159.000000 160.000000 161.000000 162.000000 163.000000 164.000000 165.000000 166.000000 167.000000 168.000000 169.000000 170.000000 171.000000 172.000000 173.000000 174.000000 175.000000 176.000000 177.000000 178.000000 179.000000 180.000000 181.000000 182.000000 183.000000 184.000000 185.000000 186.000000 187.000000 188.000000 189.000000 190.000000 191.000000 192.000000 193.000000 194.000000 195.000000 196.000000 197.000000 198.000000 199.000000
```

MPI_Allreduce

- ▶ 将MPI_Reduce的结果发送到其他所有进程可以直接通过MPI_Allreduce实现
- ▶ 即除了根进程外，所有其他进程有时也需要得到缩减后的结果



- ▶ $\text{MPI_Allreduce} = \text{MPI_Reduce} + \text{MPI_Bcast}$

MPI_Allreduce

- ▶ `int MPI_Allreduce(const void *sendbuf,
void* recvbuf, int count,
MPI_Datatype datatype,
MPI_Op op, MPI_Comm comm);`
- ▶ 和MPI_Reduce相比，不用输入root进程

MPI_Allreduce例子：计算标准差

- ▶ 进程0有数组{1,2,...,100}
- ▶ 进程1有数组{101,102,...,200}
- ▶ 进程2有数组{201,202,...,300}
- ▶ 计算合并后数组的标准差
- ▶ 利用MPI_Allreduce传递整体均值。

MPI_Allreduce例子：计算标准差

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h>

float mean(float *num, int length){
    float s = 0;
    for(int i=0; i<length; i++){
        s += num[i];
    }
    s /= length;
    return s;
}

int main(){
    MPI_Init(NULL, NULL);
    int world_rank, world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    float num[100], num_avg_per_proc, num_avg;

    for(int i=0; i<100; i++){
        num[i] = world_rank*100+i+1;
    }
    num_avg_per_proc = mean(num, 100);
```

MPI_Allreduce例子：计算标准差

```
MPI_Allreduce(&num_avg_per_proc, &num_avg, 1, MPI_FLOAT, MPI_SUM,
              MPI_COMM_WORLD);

num_avg /= 3;

float sum_sq_per_proc = 0, sum_sq;
for(int i=0; i<100; i++){
    sum_sq_per_proc += (num[i]-num_avg)*(num[i]-num_avg);
}

MPI_Reduce(&sum_sq_per_proc, &sum_sq, 1, MPI_FLOAT, MPI_SUM,
           0, MPI_COMM_WORLD);

if(world_rank == 0){
    float sd;
    sd = sqrt(sum_sq / (300-1));
    printf("The standard deviation is %f\n", sd);
}

MPI_Finalize();
```

```
}
```

MPI_Allreduce例子：计算标准差

```
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpicc -o sd_allreduce sd_allreduce.c -lm
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpirun -np 3 ./sd_allreduce
The standard deviation is 86.746758
```


- 代码中我们利用了sqrt函数来计算二次根，此函数来自math头文件，这需要在编译命令中，通过-lm来进行链接（即link math）

组和通信器管理

- ▶ 为什么要引入组和通信器？
- ▶ 在之前的内容中，我们都是用的世界通信器 MPI_COMM_WORLD。对于一些简单的应用，这已经足够了。但是当应用变得更为复杂以后，我们可能只需要一部分进程间通信。所以，我们有必要知道如何定义新的通讯器。



构建新通信器示意图



组和通信器的定义

- ▶ 一个组(group)指一个有顺序的进程的集合
- ▶ 组里的每一个进程都具有一个唯一的整数rank
 - ▶ rank值从0开始到S-1结束, 其中S指组中进程个数
 - ▶ 一个组总是和一个通信器对象相联系
- ▶ 一个通信器(communicator)包含一组可能会互相通信的进程
 - ▶ 所有MPI消息传递都必须明确一个通信器
 - ▶ 比如世界通信器MPI_COMM_WORLD
- ▶ 从程序员的角度, 组和通信器是一样的。关于组的操作主要明确哪个进程用于构造通信器

组和通信器对象的主要用途

- ▶ 根据目的，将进程编组
- ▶ 能在一个进程子集上进行集体通信操作
- ▶ 为执行用户定义的虚拟拓扑提供基础
- ▶ 提供安全的通信

编程时需要注意的地方

- ▶ 组或通信器是动态的，它们能在程序执行期间创造或破坏
- ▶ 进程可以在多于一个组或通信器中；但在每一个组或通信器中，其具有唯一的rank
- ▶ MPI提供了和组、通信器等相关的超过40个常规操作

构建新通信器的方法

▶ 方法一

- ▶ 通过MPI_Group_incl构建新的组，再用MPI_Comm_create构建基于此新组的通信器

▶ 方法二

- ▶ 利用MPI_Comm_split直接划分原通信器来构建新通信器

方法一

▶ `int MPI_Comm_group(MPI_Comm comm, MPI_Group *group);`

▶ 检索和通信器comm相关联的组

▶ `int MPI_Group_incl(MPI_Group group, int n, int *ranks,
MPI_Group *newgroup);`


▶ 在当前group中构建包含特定ranks的n个进程的新组

▶ `int MPI_Comm_create(MPI_Comm comm,
MPI_Group newgroup, MPI_Comm *newcomm);`

▶ 给原始通信器comm中的新组newgroup建立对应的通信器

方法一的例子

- ▶ 定义进程0、1、...、7
- ▶ 定义新的通信器{2,3,4,5}
- ▶ 在新通信器中的各个进程，把在世界通信器中的rank归约求和到进程2上



方法一的例子

```
#include <mpi.h>
#include <stdio.h>

int main(){
    MPI_Init(NULL, NULL);

    int world_size, world_rank, sendbuf, recvbuf;
    int new_rank;
    int ranks[4] = {2,3,4,5};

    MPI_Group world_group, new_group;
    MPI_Comm new_comm;

    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    sendbuf = world_rank;

    //extract the world group handle
    MPI_Comm_group(MPI_COMM_WORLD, &world_group);

    MPI_Group_incl(world_group, 4, ranks, &new_group);

    MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);
```


方法一的例子

```
if(new_comm != MPI_COMM_NULL){
    MPI_Comm_rank(new_comm, &new_rank);
    MPI_Reduce(&sendbuf, &recvbuf, 1, MPI_INT,
               MPI_SUM, 0, new_comm);
    printf("world rank = %d, new rank = %d\n",
           world_rank, new_rank);
}

if(new_rank == 0){
    printf("*** world rank = %d, new rank = %d, sum = %d\n",
           world_rank, new_rank, recvbuf);
}

MPI_Finalize();
}
```

方法一的例子

```
if(new_comm != MPI_COMM_NULL){  
    MPI_Comm_rank(new_comm, &new_rank);  
    MPI_Reduce(&sendbuf, &recvbuf, 1, MPI_INT,  
              MPI_SUM, 0, new_comm);  
    printf("world rank = %d, new rank = %d\n",  
          world_rank, new_rank);  
}  
  
if(new_rank == 0){  
    printf("*** world rank = %d, new rank = %d, sum = %d\n",  
          world_rank, new_rank, recvbuf);  
}  
  
MPI_Finalize();  
}
```



► 请注意！ MPI_Reduce 中 root 进程为新通信器中的 rank!

方法一的例子

```
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpicc -o group_comm_ex group_comm_ex.c
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpirun -np 8 ./group_comm_ex
world rank = 5, new rank = 3
world rank = 4, new rank = 2
world rank = 3, new rank = 1
world rank = 2, new rank = 0
*** world rank = 2, new rank = 0, sum = 14
```

方法二：利用MPI_Comm_split

Split a Large Communicator Into Smaller Communicators



方法二：利用MPI_Comm_split

- ▶ `int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm *newcomm);`
- ▶ `comm`指需要进行分裂的通信器
- ▶ `color`指相同`color`的进程在相同的通信器，`color`的值需要非负
- ▶ `key`指在新通信器中各个进程的相对`rank`
 - ▶ 对于相同`color`的两个进程，当`key`值相同时，排列根据原始通信器中的`rank`
- ▶ 返回值`newcomm`表示指向一个新通信器的句柄
- ▶ `MPI_Comm_split`构建的通信器是**不能相互重叠**的
- ▶ 如果`color`是`MPI_UNDEFINED`，对应的那个进程不被任何一个新的通讯器包含。

方法二的例子

```
#include <mpi.h>
#include <stdio.h>

int main(){
    MPI_Init(NULL, NULL);

    int world_rank, world_size;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    int color = world_rank / 4;

    MPI_Comm new_comm;
    MPI_Comm_split(MPI_COMM_WORLD, color, world_rank, &new_comm);

    int new_rank, new_size;
    MPI_Comm_rank(new_comm, &new_rank);
    MPI_Comm_size(new_comm, &new_size);

    printf("world rank/size: %d/%d --- new rank/size: %d/%d\n",
           world_rank, world_size, new_rank, new_size);


    MPI_Finalize();
}
```

方法二的例子

```
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpicc -o split split.c
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpirun -np 16 ./split
world rank/size: 0/16 --- new rank/size: 0/4
world rank/size: 1/16 --- new rank/size: 1/4
world rank/size: 8/16 --- new rank/size: 0/4
world rank/size: 12/16 --- new rank/size: 0/4
world rank/size: 10/16 --- new rank/size: 2/4
world rank/size: 2/16 --- new rank/size: 2/4
world rank/size: 9/16 --- new rank/size: 1/4
world rank/size: 3/16 --- new rank/size: 3/4
world rank/size: 11/16 --- new rank/size: 3/4
world rank/size: 4/16 --- new rank/size: 0/4
world rank/size: 13/16 --- new rank/size: 1/4
world rank/size: 14/16 --- new rank/size: 2/4
world rank/size: 5/16 --- new rank/size: 1/4
world rank/size: 15/16 --- new rank/size: 3/4
world rank/size: 6/16 --- new rank/size: 2/4
world rank/size: 7/16 --- new rank/size: 3/4
```

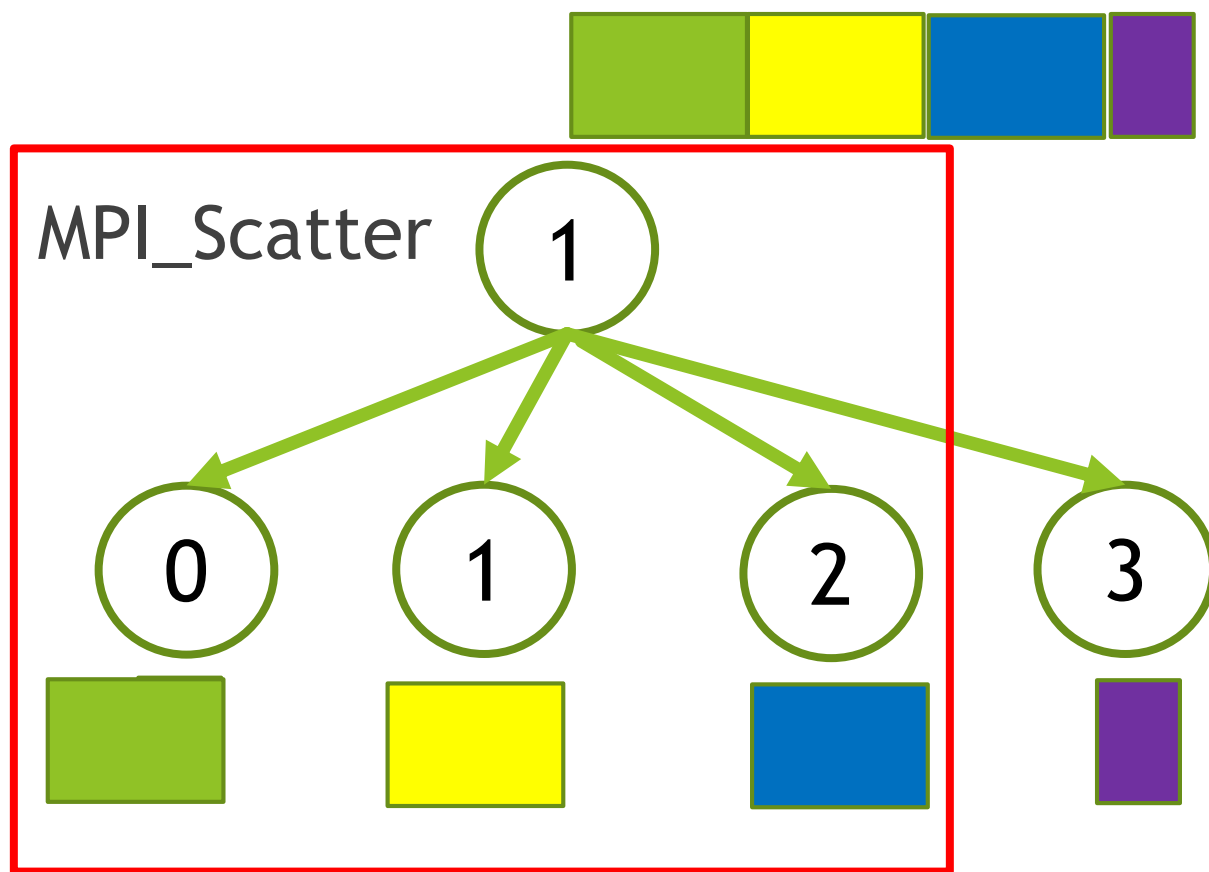
不能均分数据的例子

- ▶ MPI_Scatter中数据不可均分的例子
- ▶ 进程1的数组{1,2,3,4,5,6,7,8,9,10}分给四个进程




不能均分数据的例子

- ▶ 算法思路：先定义一个通信器 (0,1,2)，在新通信器上用 MPI_Scatter，将剩下的利用点对点通信传给进程3。



不能均分数据的例子

- ▶ 算法思路：先定义一个通信器 (0,1,2)，在新通信器上用 MPI_Scatter，将剩下的利用点对点通信传给进程3。



不能均分数据的例子

```
int main(){
    MPI_Init(NULL, NULL);

    int world_size, world_rank, *sendbuf, *recvbuf;
    int recvbuf_proc3;
    int new_rank;
    int ranks[3] = {0,1,2};

    MPI_Group world_group, new_group;
    MPI_Comm new_comm;

    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    if(world_rank == 1){
        sendbuf = (int *)malloc(10*sizeof(int));
        for(int i=0; i < 10; i++){
            sendbuf[i] = i+1;
        }
    }

    if(world_rank < 3){
        recvbuf = (int *)malloc(3*sizeof(int));
    }
}
```

不能均分数据的例子

```
//scatter numbers to processes 0-2
MPI_Comm_group(MPI_COMM_WORLD, &world_group);

MPI_Group_incl(world_group, 3, ranks, &new_group);
MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);

if(new_comm != MPI_COMM_NULL){
    MPI_Scatter(sendbuf, 3, MPI_INT, recvbuf, 3, MPI_INT,
               1, new_comm);
    printf("Process %d received numbers %d %d %d\n",
           world_rank, recvbuf[0], recvbuf[1], recvbuf[2]);
}

//send the remaining number to process 3
if(world_rank == 1){
    MPI_Send(&(sendbuf[9]), 1, MPI_INT, 3, 888, MPI_COMM_WORLD);
}

if(world_rank == 3){
    MPI_Recv(&recvbuf_proc3, 1, MPI_INT, 1, 888,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process %d received number %d\n",
           world_rank, recvbuf_proc3);
}
```

不能均分数据的例子

```
xiangyu@xiangyu-VirtualBox:~/parallel_computing_files$ mpirun -np 4 ./scatter_unbalanced  
Process 1 received numbers 4 5 6  
Process 3 received number 10  
Process 2 received numbers 7 8 9  
Process 0 received numbers 1 2 3
```