

6. MPI在一些统计方法上的应用

罗翔宇


中国人民大学统计与大数据研究院

MPI学习总览

- ▶ MPI点对点通信: MPI_Send, MPI_Recv
- ▶ MPI集体通信: MPI_Bcast, MPI_Scatter, MPI_Gather, MPI_Reduce
- ▶ MPI部分通信: 分组+集体通信
 - ▶ MPI_Comm_group, MPI_Group_incl, MPI_Comm_create
 - ▶ MPI_Comm_split


什么是Kmeans

- ▶ Kmeans是一种聚类(clustering)算法，将若干数据聚集成K类
- ▶ Kmeans不需要数据的标签，它是非监督(unsupervised)
 - ▶ 和分类(classification)不同



什么是Kmeans

- ▶ Kmeans是一种聚类(clustering)算法，将若干数据聚集成K类
- ▶ Kmeans不需要数据的标签，它是非监督(unsupervised)
 - ▶ 和分类(classification)不同



Kmeans的数学表示


- ▶ 假设我们有N个数据点 x_1, x_2, \dots, x_N ，计划用Kmeans将其聚为K类
- ▶ 假设数据点 x_i 来自于类型 $Z_i \in \{1, 2, \dots, K\}$
- ▶ 用 C_k 表示属于类型k的数据点的角标集合，即 $C_k = \{i: Z_i = k\}$
- ▶ 因此 $C_k (1 \leq k \leq K)$ 就构成对 $\{1, 2, \dots, N\}$ 的一个划分
- ▶ Kmeans的目的为优化如下目标函数

$$\min_{C_1, \dots, C_K} \sum_{1 \leq k \leq K} \sum_{i \in C_k} \left(x_i - \frac{1}{|C_k|} \sum_{i \in C_k} x_i \right)^2$$

- ▶ $d_k = \frac{1}{|C_k|} \sum_{i \in C_k} x_i$ 表示类k的中心点
- ▶ 但是，此问题是NP hard


Kmeans的迭代算法

- ▶ 迭代算法举例如下($K=3$)
- ▶ 第0步, 初始化中心点 d_1, d_2, d_3



Kmeans的迭代算法


- ▶ 第1步，计算每个点到中心点的距离，并将此点归于距离最小的那个类



Kmeans的迭代算法


► 第2步，根据每个点所在类，重新计算每个类的中心点

►
$$d_k = \frac{1}{|C_k|} \sum_{i \in C_k} x_i$$




Kmeans的迭代算法

► 不断重复第1步、第2步，直到结果稳定。




Kmeans的迭代算法

- ▶ 不断重复第1步、第2步，直到结果稳定。




Kmeans的迭代算法

- ▶ 不断重复第1步、第2步，直到结果稳定。




Kmeans的迭代算法

► 不断重复第1步、第2步，直到结果稳定。



Kmeans的迭代算法两个关键函数

- ▶ 第一步：计算 Z_i 。每个点到各个中心点的距离，找到最短的那个中心点，把自己的类型调整为最短中心点所在类型
- ▶ 第二步：计算 d_k 。根据现有每个点所在类型，更新中心点



Kmeans的迭代算法并行化

- ▶ 当数据量很大的时候，如何提升Kmeans算法的速度呢？
- ▶ 假设我们有如下数据结构

x1	x2	x3	...	xN-1	xN
0.1	0.2	-0.5	...	0.1	-0.5
-0.2	-0.2	0.3	...	-0.1	0.25


- ▶ 我们将其划分为若干部分,每一部分在对应进程上计算 Z_i

x1	x2	x3	...	xN-1	xN
0.1	0.2	-0.5	...	0.1	-0.5
-0.2	-0.2	0.3	...	-0.1	0.25

- ▶ 接着将 Z_i 的值汇集到进程0上，计算 d_k


Kmeans的迭代算法并行化

- ▶ 并行Kmeans算法示意图
- ▶ 分散数据(允许数据不均匀分散)[MPI_Scatter + MPI_Group_incl]




Kmeans的迭代算法并行化

- ▶ 分散中心点
- ▶ MPI_Bcast




Kmeans的迭代算法并行化

- ▶ 第一步：每个进程上并行计算 Z_i (每个点所在类型)并汇总于进程0
- ▶ MPI_Gather



Kmeans的迭代算法并行化

- ▶ 第二步：在进程0上更新中心点，
- ▶ 并广播到各个进程（第一次迭代结束）



Kmeans并行化代码:

```
#include <math.h>

#define N 9999999
#define K 3
#define PI 3.141592654

//sample a random number on (0,1)
double runif(){
    double tmp;
    do{
        tmp = (double) rand() / (double) RAND_MAX;
    }while(tmp>=1 || tmp <=0);
    return tmp;
}

//sample Gaussian-distributed random numbers via Box-Muller algorithm
double rnorm(double mean, double sd){
    double tmp;
    double U, V;
    U = runif();
    V = runif();

    tmp = sqrt(-2.0*log(U)) * sin(2.0*PI*V);
    tmp = mean + sd*tmp;

    return tmp;
}
```

1千万减1个点,
3类
定义从正态分
布中进行抽样
的函数

Kmeans并行化代码:

```
//define a data matrix
double** make2Darray(int a1, int a2){
    double *tmp_vec = (double *)malloc(a1*a2*sizeof(double));
    double **tmp;
    tmp = (double **)malloc(a1*sizeof(double *));
    for(int i=0; i<a1; i++){
        tmp[i] = &(tmp_vec[i*a2]);
    }
    return tmp;
}

//free a data matrix
void delet2Darray(double **tmp){
    free(tmp[0]);
    free(tmp);
}
```

定义数据
矩阵

Kmeans并行化代码:

```
//compute distances between data and centers
// and return which cluster each data point belongs to

int* compute_Z(double **dat, int n_dat, int p, double **center){
    double dist, dist_tmp;
    int *Z;
    Z = (int*) malloc(n_dat * sizeof(int));

    for(int i=0; i < n_dat; i++){
        for(int k=0; k < K; k++){
            dist_tmp = 0;
            for(int j=0; j<p; j++){
                dist_tmp += pow(dat[i][j] - center[k][j], 2);
            }
            if(k==0){
                dist = dist_tmp;
                Z[i] = 0;
            }else if(dist_tmp < dist){
                dist = dist_tmp;
                Z[i] = k;
            }
        }
    }

    return Z;
}
```

定义执行
Kmeans第
一步的函数

Kmeans并行化代码:

```
//compute the centers
double ** compute_d(double **dat, int p, int *Z){
    double s;
    int num;
    double **d = make2Darray(K, p);
    for(int k=0; k < K; k++){
        for(int j=0; j < p; j++){
            s=0;
            num=0;
            for(int i=0; i < N; i++){
                if(Z[i] == k){
                    s += dat[i][j];
                    num += 1;
                }
            }
            d[k][j] = s / num;
        }
    }
    return d;
}
```

定义执行
Kmeans第
二步的函数

Kmeans并行化代码:

```
int main(){
    MPI_Init(NULL, NULL);
    int world_rank, world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    //create a new communicator
    int *rank;
    rank = (int*) malloc(sizeof(int)*(world_size-1));
    for(int i=0; i < world_size-1; i++){
        rank[i] = i;
    }

    int *Z_partial;

    MPI_Group world_group, new_group;
    MPI_Comm new_comm;

    MPI_Comm_group(MPI_COMM_WORLD, &world_group);
    MPI_Group_incl(world_group, world_size-1, rank, &new_group);
    MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);
```

定义新的通信器

Kmeans并行化代码:

```
//generate data on process 0
double **dat, **dat_partial;
int p=2;
if(world_rank == 0){
    srand(123);
    dat = make2Darray(N, p); //rows are samples
    for(int i=0; i < N; i++){
        if(i < N/3){
            dat[i][0] = rnorm(1, 0.2);
            dat[i][1] = rnorm(1, 0.2);
        }else if(i < 2*(N/3)){
            dat[i][0] = rnorm(1, 0.2);
            dat[i][1] = rnorm(-1, 0.2);
        }else{
            dat[i][0] = rnorm(-1, 0.2);
            dat[i][1] = rnorm(-1, 0.2);
        }
    }
}
```

在进程0上生成
用于聚类的数据

请注意：行对应
样本，而非列对
应样本，为什么？

Kmeans并行化代码:

```
//scatter data
if(world_rank < world_size - 1){
    dat_partial = make2Darray(N/world_size, p);
}else{
    dat_partial = make2Darray(N - (world_size-1)*(N/world_size),
                              p);
}

if(new_comm != MPI_COMM_NULL){
    MPI_Scatter(dat[0], p*(N/world_size), MPI_DOUBLE,
               dat_partial[0], p*(N/world_size), MPI_DOUBLE,
               0, new_comm);
}

if(world_rank == 0){
    MPI_Send(dat[(world_size-1)*(N/world_size)],
             p*(N - (world_size-1)*(N/world_size)),
             MPI_DOUBLE,
             world_size-1, 888, MPI_COMM_WORLD);
}

if(world_rank == world_size -1){
    MPI_Recv(dat_partial[0], p*(N - (world_size-1)*(N/world_size)),
             MPI_DOUBLE, 0, 888,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}
```

将数据分散
出去

Kmeans并行化代码:

```
//initialize centers on process 0
double **d;
d = make2Darray(K, p);

if(world_rank == 0){
    d[0][0] = 0;
    d[0][1] = 0;

    d[1][0] = 1;
    d[1][1] = -3;

    d[2][0] = -1;
    d[2][1] = -3;
    printf("====Initialization=====\n");
    printf("cluster 1's center is (%f, %f)\n", d[0][0], d[0][1]);
    printf("cluster 2's center is (%f, %f)\n", d[1][0], d[1][1]);
    printf("cluster 3's center is (%f, %f)\n", d[2][0], d[2][1]);
}

MPI_Bcast(d[0], K*p, MPI_DOUBLE, 0, MPI_COMM_WORLD);
```

初始化中心
点

Kmeans并行化代码:

```
//Kmeans
int t, *Z;

if(world_rank == 0){
    Z = (int*)malloc(N*sizeof(int));
}

double t1, t2;

MPI_Barrier(MPI_COMM_WORLD);
if(world_rank == 0){
    t1 = MPI_Wtime();
}
for(t=0; t < 5; t++){
    //step 1: obtain the cluster indicators
    if(world_rank < world_size - 1){
        Z_partial = compute_Z(dat_partial, N/world_size, p, d);
    }else{
        Z_partial = compute_Z(dat_partial,
                               N - (world_size-1)*(N/world_size),
                               p, d);
    }
    MPI_Barrier(MPI_COMM_WORLD);
}
```

执行
Kmeans

Kmeans并行化代码:

```
//gathering
if(new_comm != MPI_COMM_NULL){
    MPI_Gather(Z_partial, N/world_size, MPI_INT, Z,
              N/world_size, MPI_INT, 0, new_comm);
}

if(world_rank == world_size-1){
    MPI_Send(Z_partial, N - (world_size-1)*(N/world_size),
             MPI_INT, 0, 888, MPI_COMM_WORLD);
}

if(world_rank == 0){
    MPI_Recv(&(Z[(world_size-1)*(N/world_size)]),
            N - (world_size-1)*(N/world_size),
            MPI_INT, world_size-1, 888, MPI_COMM_WORLD,
            MPI_STATUS_IGNORE);
}
```

执行
Kmeans

Kmeans并行化代码:

```
    //step2: compute centers on processor 0
    if(world_rank == 0){
        d = compute_d(dat, p, Z);
        printf("==== Iteration %d =====\n", t);

        printf("cluster 1's center is (%f, %f)\n", d[0][0], d[0][1]);
        printf("cluster 2's center is (%f, %f)\n", d[1][0], d[1][1]);
        printf("cluster 3's center is (%f, %f)\n", d[2][0], d[2][1]);
    }
    MPI_Bcast(d[0], K*p, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Barrier(MPI_COMM_WORLD);

}

if(world_rank == 0){
    t2 = MPI_Wtime();
    printf("\nTime cost: %f seconds \n", t2 - t1);
}

MPI_Finalize();
}
```

执行
Kmeans

运行结果：利用两个核

```
[luoxiangyu@node-1-14 parallel_computing]$ mpirun -np 2 ./Kmeans
=====Initialization=====
cluster 1's center is (0.000000, 0.000000)
cluster 2's center is (1.000000, -3.000000)
cluster 3's center is (-1.000000, -3.000000)
===== Iteration 0 =====
cluster 1's center is (0.346462, -0.291954)
cluster 2's center is (1.127639, -1.380798)
cluster 3's center is (-1.126114, -1.380772)
===== Iteration 1 =====
cluster 1's center is (0.982360, 0.878207)
cluster 2's center is (1.018051, -1.025161)
cluster 3's center is (-1.000923, -1.000578)
===== Iteration 2 =====
cluster 1's center is (1.000094, 1.000196)
cluster 2's center is (1.000029, -1.000020)
cluster 3's center is (-1.000063, -0.999942)
===== Iteration 3 =====
cluster 1's center is (1.000094, 1.000199)
cluster 2's center is (1.000028, -1.000018)
cluster 3's center is (-1.000063, -0.999942)
===== Iteration 4 =====
cluster 1's center is (1.000094, 1.000199)
cluster 2's center is (1.000028, -1.000018)
cluster 3's center is (-1.000063, -0.999942)

Time cost: 1.984286 seconds
```

运行结果：利用四个核

```
[luoxiangyu@node-1-14 parallel_computing]$ mpirun -np 4 ./Kmeans
=====Initialization=====
cluster 1's center is (0.000000, 0.000000)
cluster 2's center is (1.000000, -3.000000)
cluster 3's center is (-1.000000, -3.000000)
===== Iteration 0 =====
cluster 1's center is (0.346462, -0.291954)
cluster 2's center is (1.127639, -1.380798)
cluster 3's center is (-1.126114, -1.380772)
===== Iteration 1 =====
cluster 1's center is (0.982360, 0.878207)
cluster 2's center is (1.018051, -1.025161)
cluster 3's center is (-1.000923, -1.000578)
===== Iteration 2 =====
cluster 1's center is (1.000094, 1.000196)
cluster 2's center is (1.000029, -1.000020)
cluster 3's center is (-1.000063, -0.999942)
===== Iteration 3 =====
cluster 1's center is (1.000094, 1.000199)
cluster 2's center is (1.000028, -1.000018)
cluster 3's center is (-1.000063, -0.999942)
===== Iteration 4 =====
cluster 1's center is (1.000094, 1.000199)
cluster 2's center is (1.000028, -1.000018)
cluster 3's center is (-1.000063, -0.999942)

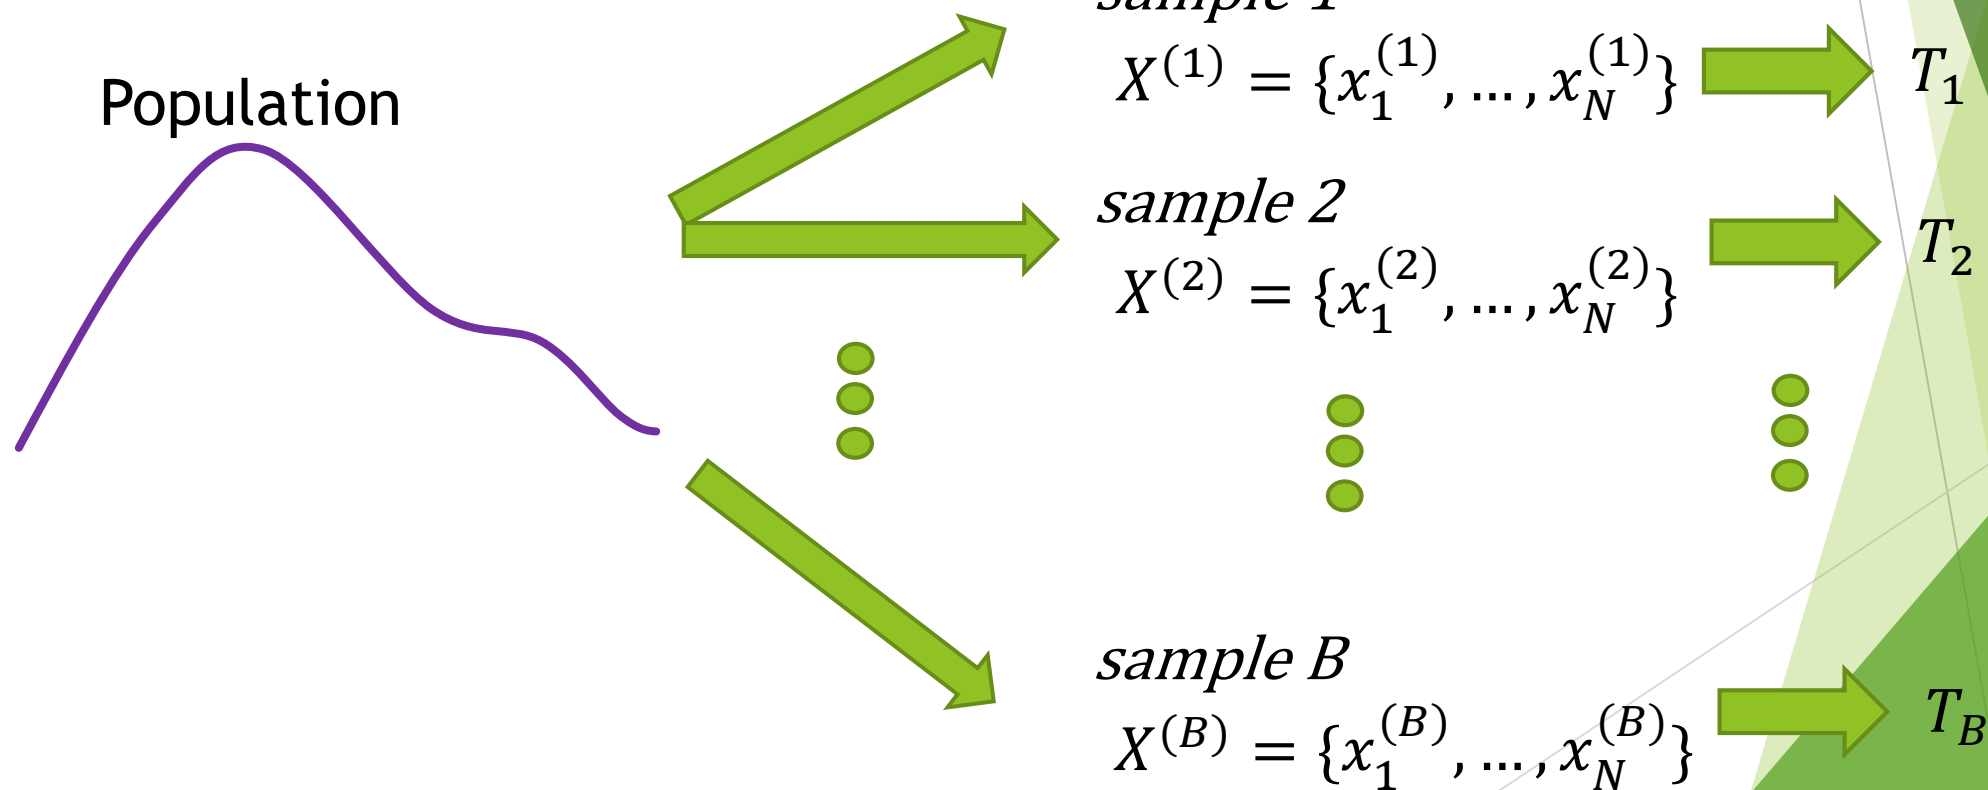
Time cost: 1.567688 seconds
```

运行结果

- ▶ 2个核：1.984秒
- ▶ 4个核：1.568秒
- ▶ 为什么速度不是一半？
 - ▶ 每次迭代，都需要进程0与其他进程进行通信(Gather, Scatter,)
 - ▶ 电脑也被其他程序所占用
 - ▶ 数据不能均匀等分，有一个额外的点对点通信
 - ▶ ...

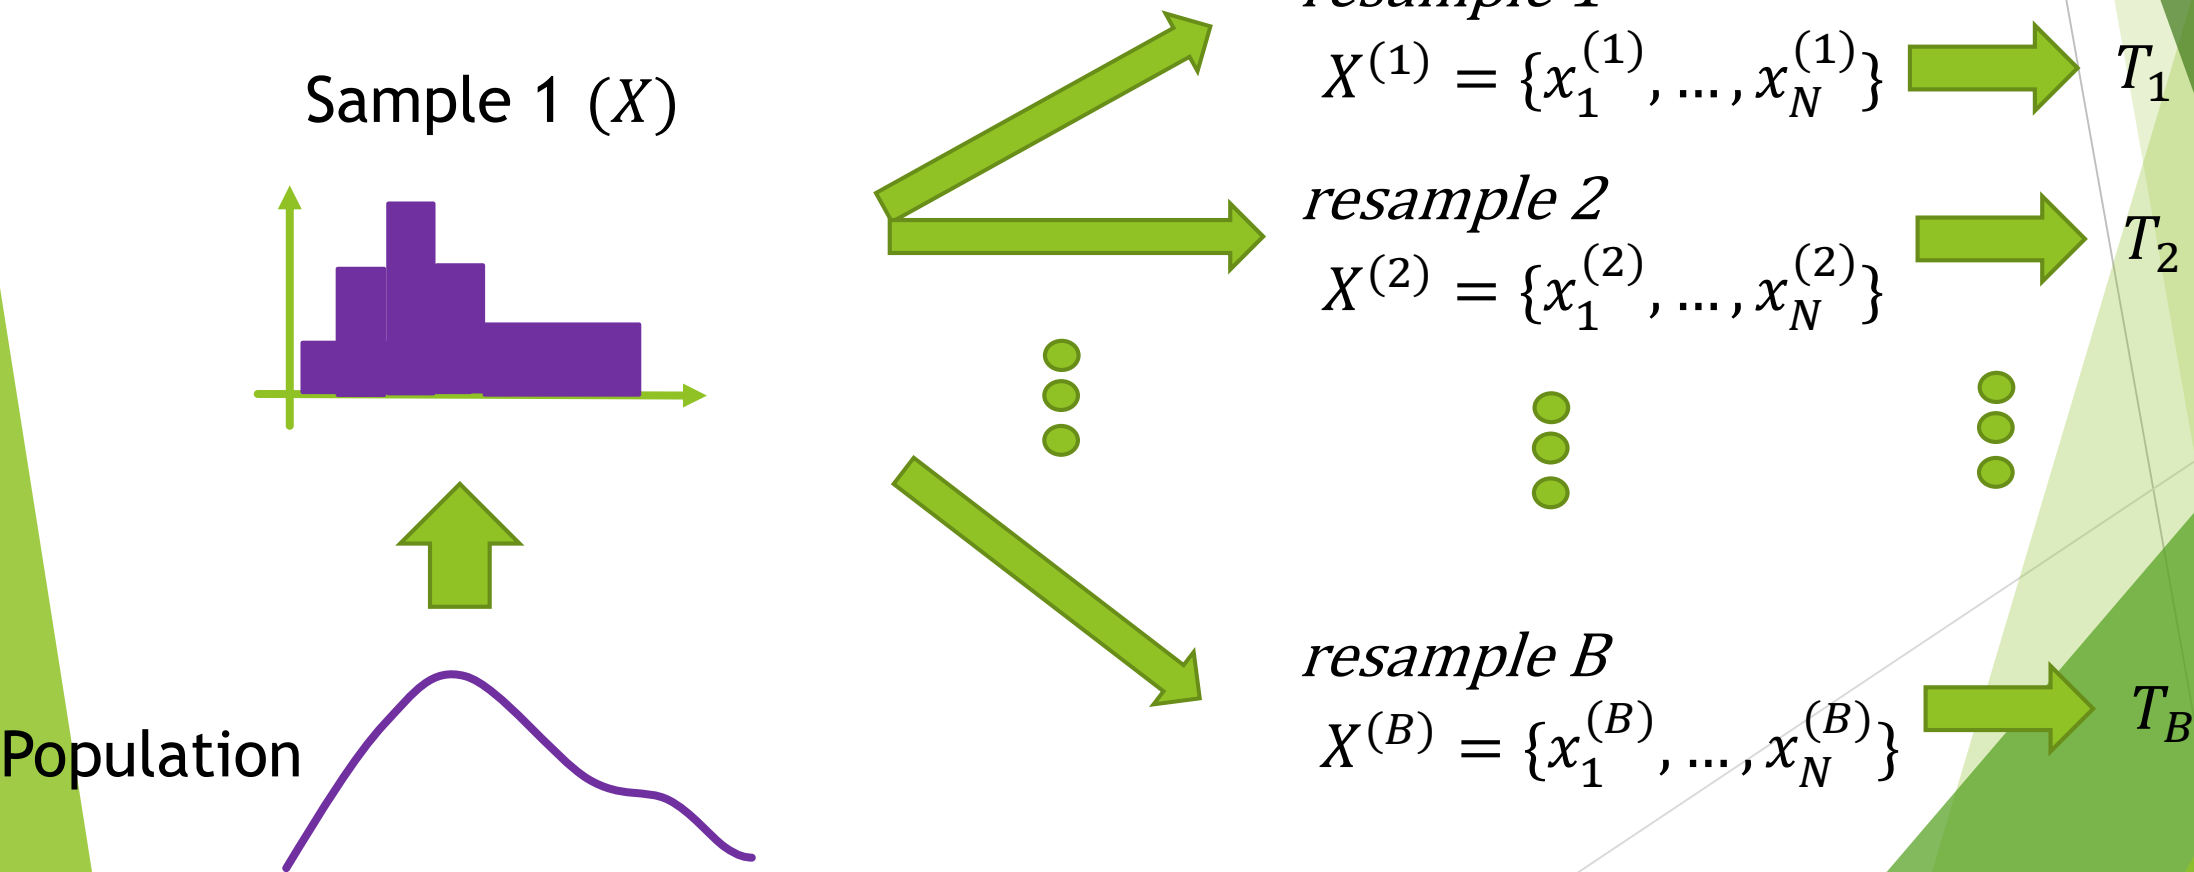
Bootstrap

- ▶ 假设我们有数据 $X = \{x_1, x_2, \dots, x_N\}$, Bootstrap 的目的是刻画出某个依赖于 X 的统计量 $T(X)$ 的分布
- ▶ 理想情形:



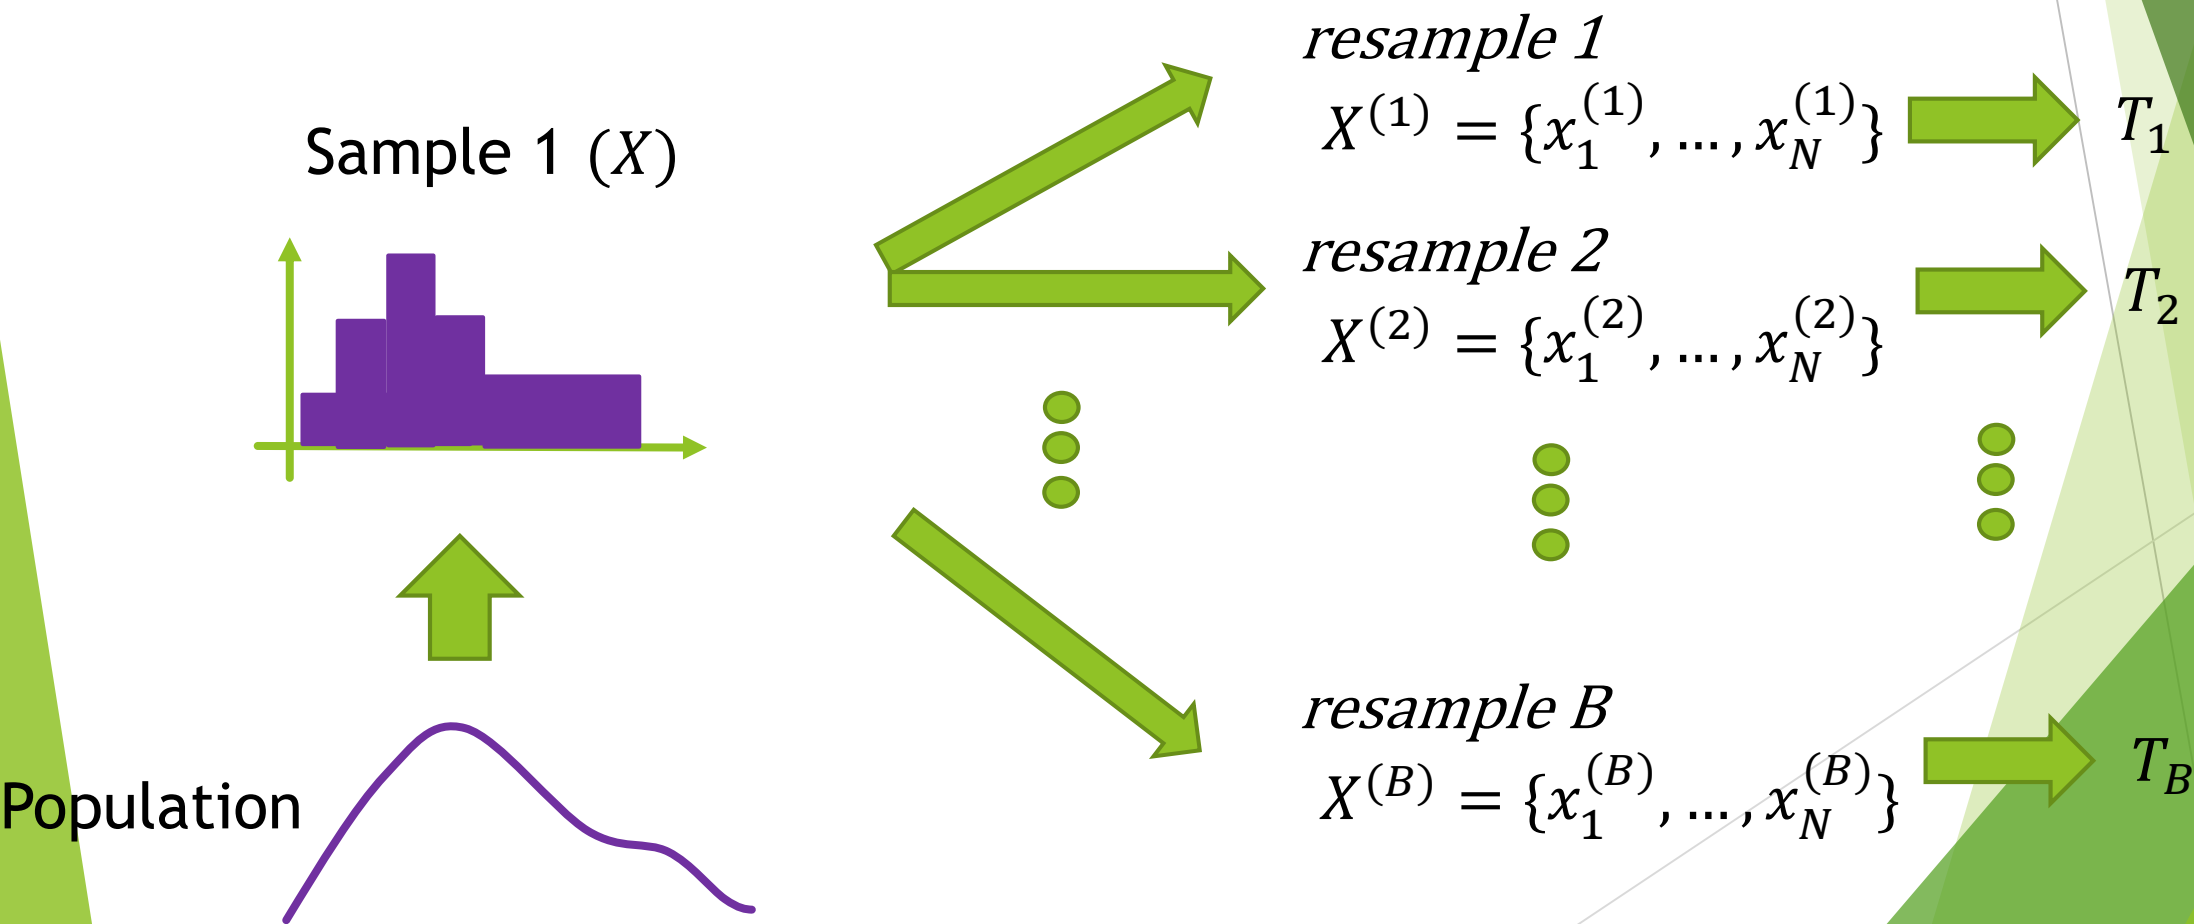
Bootstrap

- ▶ 假设我们有数据 $X = \{x_1, x_2, \dots, x_N\}$, Bootstrap 的目的是刻画出某个依赖于 X 的统计量 $T(X)$ 的分布
- ▶ 实际情形 (Bootstrap):




Bootstrap

- ▶ 重抽样的过程是**有放回地抽取**，每次之间**相互独立**
- ▶ 实际情形(Bootstrap):



Bootstrap并行化

- ▶ 例子： $X = \{x_1, x_2, \dots, x_{100}\}$ 来自于标准正态分布，通过Bootstrap求其样本方差的均值($E(S^2)$)的估计值
- ▶ 重抽样次数选为 $B=10$ 万
- ▶ 调用 S 个进程，每个进程做重抽样 B/S 次



Bootstrap并行化

```
#include <math.h>

#define B 100000
#define N 100
#define PI 3.141592654

//sample a random number on (0,1)
double runif(){
    double tmp;
    do{
        tmp = (double) rand() / (double) RAND_MAX;
    }while(tmp>=1 || tmp <=0);
    return tmp;
}

//sample Gaussian-distributed random numbers via Box-Muller algorithm
double rnorm(double mean, double sd){
    double tmp;
    double U, V;
    U = runif();
    V = runif();

    tmp = sqrt(-2.0*log(U)) * sin(2.0*PI*V);
    tmp = mean + sd*tmp;

    return tmp;
}
```

定义抽取随机数函数

Bootstrap并行化

```
//sum function
double sum(double* array, int length){
    double s = 0;
    for(int i=0; i<length; i++){
        s += array[i];
    }
    return s;
}

//var function
double var(double* array, int length){
    double avg;
    avg = sum(array, length) / length;
    double s = 0;
    for(int i=0; i<length; i++){
        s += (array[i]-avg)*(array[i]-avg);
    }
    s = s/(length - 1);
    return s;
}
```

求方差函数

Bootstrap并行化

```
//uniformly sample from 0,1,...,n-2,n-1
int sample(int n){
    double U = runif();
    int tmp;
    for(int i=0; i<n; i++){
        if((U > i/(double)n)&&(U < (i+1)/(double)n)){
            tmp = i;
            break;
        }
    }
    return tmp;
}

int main(){
    MPI_Init(NULL, NULL);
    int world_rank, world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    double *dat;
    dat = (double *)malloc(sizeof(double)*N);
```

随机抽取0
至n-1中整
数的函数

Bootstrap并行化

```
if(world_rank == 0){
    srand(123);
    for(int i=0; i<N; i++)
        dat[i] = rnorm(0,1);
}

MPI_Bcast(dat, N, MPI_DOUBLE, 0, MPI_COMM_WORLD);

int B0 = B / world_size;

//Bootstrap
int t;
double s=0;
double* re_dat = (double *)malloc(sizeof(double)*N);
int index;

double t1, t2;
MPI_Barrier(MPI_COMM_WORLD);
if(world_rank == 0){
    t1 = MPI_Wtime();
}
```

定义并广播
数组

Bootstrap并行化

```
for(t=0; t<B0; t++){
    for(int i=0; i<N; i++){
        index = sample(N);
        re_dat[i] = dat[index];
    }
    s += var(re_dat, N);
}

MPI_Barrier(MPI_COMM_WORLD);
if(world_rank == 0){
    t2 = MPI_Wtime();

double avg;
MPI_Reduce(&s, &avg, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if(world_rank == 0){
    avg = avg / B;
    printf("The bootstrap estimate of the mean of the sample variance is %f\n", avg);
    printf("Time cost %f seconds\n", t2 - t1);
}

MPI_Finalize();
}
```

进行
Bootstrap

Bootstrap并行化结果

```
[luoxiangyu@node-3-2 parallel_computing]$ mpirun -np 2 ./Bootstrap
The bootstrap estimate of the mean of the sample variance is 1.207092
Time cost 4.004207 seconds
[luoxiangyu@node-3-2 parallel_computing]$ mpirun -np 4 ./Bootstrap
The bootstrap estimate of the mean of the sample variance is 1.207495
Time cost 2.002266 seconds
[luoxiangyu@node-3-2 parallel_computing]$ mpirun -np 8 ./Bootstrap
The bootstrap estimate of the mean of the sample variance is 1.205771
Time cost 1.001832 seconds
```


可以看到Bootstrap并行中除了开始和结束，

中间过程不需要任何通信开销，

因此增加一倍的进程(核数允许的情况下)，时间缩小一倍

矩阵乘法的并行化

- ▶ 矩阵乘法，假设有两个方阵A和B，想计算A和B的矩阵乘积



$$(AB)_{ij} = \sum_{1 \leq k \leq n} A_{ik} B_{kj}$$